Graph Contrastive Learning with Augmentations
(Appendix)

Yuning You[1*], Tianlong Chen[2*], Yongduo Sui[3], Ting Chen[4], Zhangyang Wang[2], Yang Shen[1]
[1]Texas A&M University, [2]University of Texas at Austin,
[3]University of Science and Technology of China, [4]Google Research, Brain Team
{yuning.you,yshen}@tamu.edu, {tianlong.chen,atlaswang}@utexas.edu
syd2019@mail.ustc.edu.cn, iamtingchen@google.com

A Contrastive Learning and Augmentation Algorithms

Algorithm 1 Graph Contrastive Learning

Initialize: Data $\{\mathcal{G}_i : i \in M\}$, $f(\cdot)$, $g(\cdot)$, $\mathcal{T}$
1: **for** sampled minibatch of data $\{\mathcal{G}_i : i \in N\}$ **do**
2:     **for** $k = 1$ **to** N **do**
3:         Sample q_i, q_j from $\mathcal{T}$
4:         $\hat{\mathcal{G}}_{2k-1} \sim q_1(\mathcal{G}_k|\mathcal{G})$ # 1$^{\text{st}}$ augmentation
5:         $h_{2k-1} = f(\hat{\mathcal{G}}_{2k-1})$, $z_{2k-1} = g(h_{2k-1})$
6:         $\hat{\mathcal{G}}_{2k} \sim q_2(\mathcal{G}_k|\mathcal{G})$ # 2$^{\text{nd}}$ augmentation
7:         $h_{2k} = f(\hat{\mathcal{G}}_{2k})$, $z_{2k} = g(h_{2k})$
8:     **end for**
9:     **for** $\forall i \in \{1, 3, 5..., 2N-1\}$ and $\forall j \in \{2, 4, 6, ..., 2N\}$ **do**
10:         $s_{i,j} = z_i^{\mathsf{T}} z_j / (\tau \|z_i\| \|z_j\|)$
11:     **end for**
12:     **define** $\ell(i,j) = -\log \frac{\exp(s_{i,j})}{\sum_{k=1}^{N} \mathbb{1}_{[2k \neq i+1]} \exp(s_{i,2k})}$
13:
14:     $\mathcal{L} = \frac{1}{N} \sum_{k=1}^{N} \ell(2k-1, 2k)$
15:     # Asymmetric and simplified compared to the SimCLR loss
16:     Update encoder $f(\cdot)$ and $g(\cdot)$ to minimize $\mathcal{L}$
17: # maximize agreement
18: **end for**
19: **return** Encoder $f(\cdot)$

Algorithm 2 Data Augmentation: Subgraph

Initialize: Graph $\mathcal{G} = \{\mathcal{V}, \mathcal{E}\}$, augmentation ratio k, sampled graph $\mathcal{G}_{\text{samp}} = \{\mathcal{V}_{\text{samp}}, \mathcal{E}_{\text{samp}}\}$ where $\mathcal{V}_{\text{samp}} = \mathcal{E}_{\text{samp}} = \emptyset$, neighbor vertex set $\mathcal{V}_{\text{neigh}} = \emptyset$
1: Sample a node $v \in \mathcal{V}$, such that $\mathcal{V}_{\text{samp}} = \{v\}$ and $\mathcal{V}_{\text{neigh}} = \mathcal{N}(v)$
2: **while** $|\mathcal{V}_{\text{samp}}| \leq k|\mathcal{V}|$ **do**
3:      Sample a node $v \in \mathcal{V}_{\text{neigh}}$
4:      **if** $v \in \mathcal{V}_{\text{samp}}$ **then**
5:          Continue
6:      **end if**
7:      Update $\mathcal{V}_{\text{samp}} = \mathcal{V}_{\text{samp}} \cup \{v\}$, $\mathcal{V}_{\text{neigh}} = \mathcal{V}_{\text{neigh}} \cup \mathcal{N}(v)$
8: **end while**
9: Update $\mathcal{E}_{\text{samp}} = \{e | e \in \mathcal{E} \text{ and } (e[0] \in \mathcal{V}_{\text{samp}} \text{ or } e[1] \in \mathcal{V}_{\text{samp}})\}$
10: **return** $\mathcal{G}_{\text{samp}}$

Algorithm 3 Data Augmentation: Subgraph-W

Initialize: Graph $\mathcal{G} = \{\mathcal{V}, \mathcal{E}\}$, augmentation ratio k, sampled graph $\mathcal{G}_{\text{samp}} = \{\mathcal{V}_{\text{samp}}, \mathcal{E}_{\text{samp}}\}$ where $\mathcal{V}_{\text{samp}} = \mathcal{E}_{\text{samp}} = \emptyset$, neighbor vertex set $\mathcal{V}_{\text{neigh}} = \emptyset$
1: Sample a node $v \in \mathcal{V}$, such that $\mathcal{V}_{\text{samp}} = \{v\}$ and $\mathcal{V}_{\text{neigh}} = \mathcal{N}(v)$
2: **while** $|\mathcal{V}_{\text{samp}}| \leq k|\mathcal{V}|$ **do**
3:      Update $\mathcal{V}_{\text{samp}} = \mathcal{V}_{\text{samp}} \cup \mathcal{V}_{\text{neigh}}$, $\mathcal{V}_{\text{neigh}} = \cup_{v \in \mathcal{V}_{\text{samp}}} \mathcal{N}(v)$
4: **end while**
5: Update $\mathcal{E}_{\text{samp}} = \{e | e \in \mathcal{E} \text{ and } (e[0] \in \mathcal{V}_{\text{samp}} \text{ or } e[1] \in \mathcal{V}_{\text{samp}})\}$
6: **return** $\mathcal{G}_{\text{samp}}$

Algorithm 4 Data Augmentation: Subgraph-D

Initialize: Graph $\mathcal{G} = \{\mathcal{V}, \mathcal{E}\}$, augmentation ratio k, sampled graph $\mathcal{G}_{\text{samp}} = \{\mathcal{V}_{\text{samp}}, \mathcal{E}_{\text{samp}}\}$ where $\mathcal{V}_{\text{samp}} = \mathcal{E}_{\text{samp}} = \emptyset$, neighbor vertex set $\mathcal{V}_{\text{neigh}} = \emptyset$
1: Sample a node $v \in \mathcal{V}$, such that $\mathcal{V}_{\text{samp}} = \{v\}$ and $\mathcal{V}_{\text{neigh}} = \mathcal{N}(v)$
2: **while** $|\mathcal{V}_{\text{samp}}| \leq k|\mathcal{V}|$ **do**
3:      Sample a node $v \in \mathcal{V}_{\text{neigh}}$
4:      **if** $v \in \mathcal{V}_{\text{samp}}$ **then**
5:          Continue
6:      **end if**
7:      Update $\mathcal{V}_{\text{samp}} = \mathcal{V}_{\text{samp}} \cup \{v\}$, $\mathcal{V}_{\text{neigh}} = \mathcal{N}(v)$
8: **end while**
9: Update $\mathcal{E}_{\text{samp}} = \{e | e \in \mathcal{E} \text{ and } (e[0] \in \mathcal{V}_{\text{samp}} \text{ or } e[1] \in \mathcal{V}_{\text{samp}})\}$
10: **return** $\mathcal{G}_{\text{samp}}$

B   Detailed Settings for Augmentation Experiments (Section 4 in Main Text)

We evaluate our proposed framework with different augmentation pairs in the semi-supervised learning setting on graph classification [1] via pre-training & finetuing where pre-training is performed with 100 epochs, 0.001 learning rate, and finetuning follows the 10-fold evaluation finetuning in [2] that achieves the comparable SOTA performance in the fully-supervised setting. Graph convolutional network (GCN) is adopted as the GNN-based encoder also following [2]. Experiments are performed with 1% (if there are over 10 samples for each class) and 10% label rate for 5 times with mean and standard deviation of accuracies (%) reported.

C   Graph Contrastive Learning for Superpixel Graphs

Superpixel graphs (statistics in Table S1) gain from all augmentations except attribute masking as shown in Figure S1. For node dropping, it corresponds to pixel discarding and for subgraph to

Table S1: Superpixel graph dataset statistics.

Datasets	Category	Graph Num.	Avg. Node	Avg. Degree
MNIST	Superpixel Graphs	70000	70.57	8

Figure S1: Semi-supervised learning accuracy gain (%) when contrasting different augmentation pairs, compared to training from scratch under MNIST. Pairing "Identical" stands for a no-augmentation baseline for contrastive learning, where the positive pair diminishes and the negative pair consists of two non-augmented graphs. Warmer colors indicate better performance gains. The baseline training-from-scratch accuracy is 79.71%.

cropping, which are already shown as useful augmentations in images [3]. Surprisingly, attribute masking corresponding to image completion hurts the performance, which might result from our implementation: node attributes of superpixel graphs contain information of pixel value and location, and we might only mask the pixel value part rather than all analog to image completion. We do not find a related augmentation with edge perturbation and leave it for future work.

D Difficulty of Contrastive Tasks v.s. Semi-Supervised Performance

We first note that, for edge perturbation, attribute masking, and node dropping, their extents could be an indicator of the difficulty for corresponding contrastive tasks. We observed earlier in Figure 4 (main text) that, at least for COLLAB, properly increasing the extents could enhance the downstream performances.

We also note that, for attribute masking and node dropping, their patterns could correlate the difficulty as well. With larger control factor α in masking/dropping distribution, the vertices with more connections are masked/dropped with higher probability, intuitively leading to a "harder" task. We again observed in Figure 5 (main text) that at least COLLAB performances benefited from the harder task with the pattern change, while overly simple contrastive tasks with very negative α would not help.

Table S2: Performance on contrastive learning with different implemented subgraph. The intuitively simplest subgraph-W performs the worst among the three.

Augmentations	Subgraph-W	Subgraph	Subgraph-D
PROTEINS	71.50±0.85	72.67±0.60	72.74±0.56
COLLAB	57.66±1.64	63.63±1.20	65.47±1.43

For subgraph, we propose the following variants with difficulty levels. Contrastive learning with subgraphs sampled via depth-first-search (DFS) encouraged random walk is more difficult than that via width-first-search (WFS) encouraged, since the latter preserves more structure information

(connections) to assist GNNs to recover semantic information [4]. Notice that our default subgraph encourages neither DFS nor WFS, and therefore we additionally proposed subgraph-D(FS) and subgraph-W(FS) (Algorithms summarized in Appendix A) with the intuitive difficulty rank: subgraph-W < subgraph < subgraph-D. Experiments on PROTEINS and COLLAB in Table S2 agrees with our previous conjecture, that the simplest subgraph-W yields the worst performance among the three.

E   Datasets and Training in Various Settings (Section 5 in Main Text))

Semi-supervised Learning

Table S3: Datasets statistics for semi-supervised learning and unsupervised representation learning.

Datasets	Category	Graph Num.	Avg. Node	Avg. Degree
NCI1	Biochemical Molecules	4110	29.87	1.08
PROTEINS	Biochemical Molecules	1113	39.06	1.86
DD	Biochemical Molecules	1178	284.32	715.66
MUTAG	Biochemical Molecules	188	17.93	19.79
COLLAB	Social Networks	5000	74.49	32.99
RDT-B	Social Networks	2000	429.63	1.15
RDB-M	Social Networks	2000	429.63	497.75
GITHUB	Social Networks	4999	508.52	594.87
IMDB-B	Social Networks	1000	19.77	96.53
MNIST	Superpixel Graphs	70000	70.57	8
CIFAR10	Superpixel Graphs	60000	117.63	8

For all datasets we perform experiments with 1% (if there are over 10 samples for each class) and 10% label rate for 5 times, each of which corresponds to a 10-fold evaluation as [2], with mean and standard deviation of accuracies (%) reported. For pre-training, learning rate is tuned in {0.01, 0.001, 0.0001} and epoch number in {20, 40, 60, 80, 100} where grid search is performed. We follow the default setting in [2] for finetuning that achieves the SOTA performance in the fully-supervised setting.

Unsupervised Representation Learning

Experiments are performed for 5 times each of which corresponds to a 10-fold evaluation as [5], with mean and standard deviation of accuracies (%) reported.

Transfer Learning

Table S4: Datasets statistics for transfer learning.

Datasets	Category	Utilization	Graph Num.	Avg. Node	Avg. Degree
ZINC-2M	Biochemical Molecules	Pre-Training	2000000	26.62	57.72
PPI-306K	Protein-Protein Intersection Networks	Pre-Training	306925	39.82	729.62
BBBP	Biochemical Molecules	Finetuning	2039	24.06	51.90
Tox21	Biochemical Molecules	Finetuning	7831	18.57	38.58
ToxCast	Biochemical Molecules	Finetuning	8576	18.78	38.52
SIDER	Biochemical Molecules	Finetuning	1427	33.64	70.71
ClinTox	Biochemical Molecules	Finetuning	1477	26.15	55.76
MUV	Biochemical Molecules	Finetuning	93087	24.23	52.55
HIV	Biochemical Molecules	Finetuning	41127	25.51	54.93
BACE	Biochemical Molecules	Finetuning	1513	34.08	73.71
PPI	Protein-Protein Intersection Networks	Finetuning	88000	49.35	890.77

F   Theoretical Justification

Mutual information maximization. We first conceptually depict the essence of our framework, rigorously showing that GraphCL can be viewed as a kind of mutual information maximization

between the latent representations of two kinds of augmented graphs. We rewrite GraphCL loss for each data point as:

$$\ell_{i,j} = -\log\frac{\exp(\mathrm{sim}(\boldsymbol{z}_i, \boldsymbol{z}_j)/\tau)}{\sum_{k=1}^{N}\exp(\mathrm{sim}(\boldsymbol{z}_i, \boldsymbol{z}_k)/\tau)}, \tag{1}$$

which can be rewrited for a batch of graphs as:

$$\begin{aligned}
\ell &= -\frac{1}{N}\sum_{n=1}^{N}\log\frac{\exp(\mathrm{sim}(\boldsymbol{z}_i^{(n)}, \boldsymbol{z}_j^{(n)})/\tau)}{\sum_{m=1}^{N}\exp(\mathrm{sim}(\boldsymbol{z}_i^{(n)}, \boldsymbol{z}_j^{(m)})/\tau)} \\
&= -\frac{1}{N}\sum_{n=1}^{N}[\mathrm{sim}(\boldsymbol{z}_i^{(n)}, \boldsymbol{z}_j^{(n)})/\tau - \log(\sum_{m=1}^{N}\exp(\mathrm{sim}(\boldsymbol{z}_i^{(n)}, \boldsymbol{z}_j^{(m)})/\tau))] \\
&= -\frac{1}{N}\sum_{n=1}^{N}\frac{\mathrm{sim}(g(f(\hat{\mathcal{G}}_i^{(n)})), g(f(\hat{\mathcal{G}}_j^{(n)})))}{\tau} + \frac{1}{N}\sum_{n=1}^{N}\log(\sum_{m=1}^{N}\exp(\frac{\mathrm{sim}(g(f(\hat{\mathcal{G}}_i^{(n)})), g(f(\hat{\mathcal{G}}_j^{(m)})))}{\tau})),
\end{aligned} \tag{2}$$

where $\boldsymbol{z}_i^{(n)} = g(\boldsymbol{h}_i^{(n)})$, $g(\cdot)$ is the projection head, $\boldsymbol{h}_i^{(n)} = f(\hat{\mathcal{G}}_i^{(n)})$, $f(\cdot)$ is GNN, $\hat{\mathcal{G}}_i^{(n)} \sim \mathbb{P}_{\hat{\mathcal{G}}_i}$ is the augmented graph. We rewrite (2) as the expectation form:

$$\begin{aligned}
\ell &= -\mathbb{E}_{\mathbb{P}_{(\hat{\mathcal{G}}_i, \hat{\mathcal{G}}_j)}}\frac{\mathrm{sim}(g(f(\hat{\mathcal{G}}_i)), g(f(\hat{\mathcal{G}}_j)))}{\tau} + \mathbb{E}_{\mathbb{P}_{\hat{\mathcal{G}}_i}}\log(\mathbb{E}_{\mathbb{P}_{\hat{\mathcal{G}}_j}}\exp(\frac{\mathrm{sim}(g(f(\hat{\mathcal{G}}_i)), g(f(\hat{\mathcal{G}}_j)))}{\tau})) - \log N \\
&= \mathbb{E}_{\mathbb{P}_{\hat{\mathcal{G}}_i}}\{-\mathbb{E}_{\mathbb{P}_{(\hat{\mathcal{G}}_j|\hat{\mathcal{G}}_i)}}\frac{\mathrm{sim}(g(f(\hat{\mathcal{G}}_i)), g(f(\hat{\mathcal{G}}_j)))}{\tau} + \log(\mathbb{E}_{\mathbb{P}_{\hat{\mathcal{G}}_j}}\exp(\frac{\mathrm{sim}(g(f(\hat{\mathcal{G}}_i)), g(f(\hat{\mathcal{G}}_j)))}{\tau}))\} - \log N \\
&= \mathbb{E}_{\mathbb{P}_{\hat{\mathcal{G}}_i}}\{-\mathbb{E}_{\mathbb{P}_{(\hat{\mathcal{G}}_j|\hat{\mathcal{G}}_i)}}T(f(\hat{\mathcal{G}}_i), f(\hat{\mathcal{G}}_j)) + \log(\mathbb{E}_{\mathbb{P}_{\hat{\mathcal{G}}_j}}e^{T(f(\hat{\mathcal{G}}_i), f(\hat{\mathcal{G}}_j))})\} - \log N,
\end{aligned} \tag{3}$$

where $\mathbb{P}_{(\hat{\mathbb{G}}_i, \hat{\mathbb{G}}_j)}, \mathbb{P}_{(\hat{\mathbb{G}}_j|\hat{\mathbb{G}}_i)}, \mathbb{P}_{\hat{\mathbb{G}}_i}$ are respectively the joint, conditional and marginal distribution of augmented graphs, and $T(\cdot, \cdot)$ is a learnable score function that we parametrize with the similarity function $\mathrm{sim}(\cdot, \cdot)$, temperature factor τ and the projection head $g(\cdot)$. Thus, (3) fits the formulation of the InfoNCE loss [6, 7] such that minimizing (3) is equivalent to *maximizing a lower bound of the mutual information* between the latent representations of two views of graphs $\boldsymbol{h}_i = f(\hat{\mathcal{G}}_i), \boldsymbol{h}_j = f(\hat{\mathcal{G}}_j)$. We would like to highlight the crucial role of the projection head in the framework, that provides the learnable weights to construct a function space for the score function, to reach a much tighter lower bound compared with dropping the projection head (the key role of it is also empirically verified in [3]).

General framework. Furthermore, we draw the connection between GraphCL and recently proposed contrastive learning methods that we demonstrate that GraphCL can be rewrited as a general framework unifying a broad family of contrastive learning methods on graph-structured data, through our rewriting (3) as (we neglect $\log N$ for simplicity):

$$\ell = \mathbb{E}_{\mathbb{P}_{\hat{\mathcal{G}}_i}}\{-\mathbb{E}_{\mathbb{P}_{(\hat{\mathcal{G}}_j|\hat{\mathcal{G}}_i)}}T(f_1(\hat{\mathcal{G}}_i), f_2(\hat{\mathcal{G}}_j)) + \log(\mathbb{E}_{\mathbb{P}_{\hat{\mathcal{G}}_j}}e^{T(f_1(\hat{\mathcal{G}}_i), f_2(\hat{\mathcal{G}}_j))})\}, \tag{4}$$

where we maximize a lower bound of the mutual information between $\boldsymbol{h}_i = f_1(\hat{\mathcal{G}}_i), \boldsymbol{h}_j = f_2(\hat{\mathcal{G}}_j)$ that the compositions of $(f_1, \hat{\mathcal{G}}_i), (f_2, \hat{\mathcal{G}}_j)$ determine our desired views of graphs. In our implementation, we choose $f_1 = f_2$ and generate $\hat{\mathcal{G}}_i, \hat{\mathcal{G}}_j$ through data augmentation, while with various choices of the compositions result in (4) instantiating as other specific contrastive learning algorithms including [8, 9, 10, 5, 11, 12, 13].

- **DGI, HDGI, DMGI** [8, 9, 10]. DGI intends to maximize the agreement between the local and global representations. Thus, it sets f_1 a GNN encoder, f_2 is the concatenation of f_1 and a node pooling layer, and $\hat{\mathcal{G}}_i = \hat{\mathcal{G}}_j = \mathcal{G}$. HDGI is an application of DGI in heterogenous graphs that the GNN encoder f_1 is a heterogenous GNN. DMGI is an extension of DGI into multiplex networks that perform DGI with multiple relationship types and the final representations are aggregated with multiple learned features.

- **InfoGraph** [5]. InfoGraph is an extension of DGI in graph-level representation learning, which is aimed at optimize the similarity between node embeddings and graph embeddings, and there for it sets $f_1 = f_2$ as a GNN encoder, $\hat{\mathcal{G}}_i = \mathcal{G}$ is the original graph and $\hat{\mathcal{G}}_j$ is the sampled subgraph.

- **GMI** [11]. GMI intends to maximize the agreement between the raw node & edge features and the encoded node & edge features, and therefore in the node loss ℓ_{node} it sets $\hat{\mathcal{G}}_{i,\text{node}} = \hat{\mathcal{G}}_{j,\text{node}} = \mathcal{G}$, $f_{1,\text{node}}(\cdot) = I(\cdot)$ is the identical function and $f_{2,\text{node}}(\cdot)$ is a GNN encoder, and similar for the edge loss ℓ_{edge} where $f_{2,\text{node}}(\cdot)$, $f_{2,\text{edge}}(\cdot)$ share weights. The final loss is expressed as $\ell = \ell_{\text{node}} + \ell_{\text{edge}}$.

G Experiments on Node Classification

GraphCL in node classification. GraphCL is also evaluated in the unsupervised representation learning in node classification following [8], where unsupervised methods generate node embeddings with GAT as the GNN-based encoder that are fed into a down-stream classifier as shown in Table S5, and semi-supervised learning following [14] with GCN, GIN and GAT as encoders as shown in Table S6, verifying the advantage of GraphCL.

Table S5: Comparing classification accuracy on top of learned node representations. The compared deep graph infomax (DGI, [8]) performance is from the original paper under the same experiment setting.

Methods	Cora	Citeseer
DGI	82.30±0.60	71.80±0.70
NodeDrop v.s. Identical	82.41±0.10	72.22±0.18
NodeDrop v.s. NodeDrop	81.76±0.17	73.14±0.15
EdgePert v.s. Identical	82.45±0.11	72.23±0.17
EdgePert v.s. EdgePert	82.32±0.15	73.11±0.19
AttrMask v.s. Identical	82.45±0.12	72.31±0.13
AttrMask v.s. AttrMask	81.78±0.17	72.05±0.22
Subgraph v.s. Identical	82.49±0.12	72.33±0.18
Subgraph v.s. Subgraph	81.71±0.14	73.12±0.17

Table S6: Node classification experiments with different models contrasting identical vs. different augmentation. Performance on the standard test sets of PATTERN SBM graphs. Results are averaged over 4 runs with 4 different seeds.

Models	10% baseline	NodeDrop	EdgePert	AttrMask	Subgraph
GCN	70.53±0.72	67.91±0.38	68.61±0.69	68.11±0.51	67.79±0.33
GIN	96.61±2.77	97.52±0.57	98.09±0.49	97.23±0.42	98.35±0.49
GAT	76.71±8.27	75.86±3.36	79.73±7.08	84.72±2.00	85.30±2.31

References

[1] Keyulu Xu, Weihua Hu, Jure Leskovec, and Stefanie Jegelka. How powerful are graph neural networks? arXiv preprint arXiv:1810.00826, 2018.

[2] Ting Chen, Song Bian, and Yizhou Sun. Are powerful graph neural nets necessary? a dissection on graph classification. arXiv preprint arXiv:1905.04579, 2019.

[3] Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. A simple framework for contrastive learning of visual representations. arXiv preprint arXiv:2002.05709, 2020.

[4] Aditya Grover and Jure Leskovec. node2vec: Scalable feature learning for networks. In Proceedings of the 22nd ACM SIGKDD international conference on Knowledge discovery and data mining, pages 855–864, 2016.

[5] Fan-Yun Sun, Jordan Hoffmann, and Jian Tang. Infograph: Unsupervised and semi-supervised graph-level representation learning via mutual information maximization. arXiv preprint arXiv:1908.01000, 2019.

[6] Aaron van den Oord, Yazhe Li, and Oriol Vinyals. Representation learning with contrastive predictive coding. arXiv preprint arXiv:1807.03748, 2018.

[7] Michael Tschannen, Josip Djolonga, Paul K Rubenstein, Sylvain Gelly, and Mario Lucic. On mutual information maximization for representation learning. arXiv preprint arXiv:1907.13625, 2019.

[8] Petar Velickovic, William Fedus, William L Hamilton, Pietro Liò, Yoshua Bengio, and R Devon Hjelm. Deep graph infomax. In ICLR (Poster), 2019.

[9] Yuxiang Ren, Bo Liu, Chao Huang, Peng Dai, Liefeng Bo, and Jiawei Zhang. Heterogeneous deep graph infomax. arXiv preprint arXiv:1911.08538, 2019.

[10] Chanyoung Park, Donghyun Kim, Jiawei Han, and Hwanjo Yu. Unsupervised attributed multiplex network embedding. In AAAI, pages 5371–5378, 2020.

[11] Zhen Peng, Wenbing Huang, Minnan Luo, Qinghua Zheng, Yu Rong, Tingyang Xu, and Junzhou Huang. Graph representation learning via graphical mutual information maximization. In Proceedings of The Web Conference 2020, pages 259–270, 2020.

[12] Kaveh Hassani and Amir Hosein Khasahmadi. Contrastive multi-view representation learning on graphs. arXiv preprint arXiv:2006.05582, 2020.

[13] Jiezhong Qiu, Qibin Chen, Yuxiao Dong, Jing Zhang, Hongxia Yang, Ming Ding, Kuansan Wang, and Jie Tang. Gcc: Graph contrastive coding for graph neural network pre-training. In Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining, pages 1150–1160, 2020.

[14] Vijay Prakash Dwivedi, Chaitanya K Joshi, Thomas Laurent, Yoshua Bengio, and Xavier Bresson. Benchmarking graph neural networks. arXiv preprint arXiv:2003.00982, 2020.